



How to score your agent

Avreo take-home • Self-evaluation rubric

The headline number

Run your agent across all 20 messages in `05_Inbound_Messages.json`. Compare each output to the corresponding entry in `06_Benchmark.json`. Calculate two scores:

Strict accuracy = the percentage of messages where your agent matched the benchmark on all four hard fields: *category*, *priority*, *route_to*, *needs_human_review*.

Per-field accuracy = the percentage on each field individually. This tells us where the agent is strong and where it slips.

Per-field scoring

Field	How to score	Notes
category	Exact match (1) or no match (0)	Six possible values.
priority	Exact match (1) or no match (0)	P1 / P2 / P3.
route_to	Exact match (1). Partial credit (0.5) if the primary team is right but cc is missed.	Use your own judgement on partial credit.
needs_human_review	Exact match (1) or no match (0)	Boolean.

Qualitative checks (not scored, but reported)

For each message, also assess:

Draft reply quality — does the draft hit the points listed in *draft_reply_must_include* and avoid the items in *draft_reply_must_not_include*? Use your own judgement, not a checklist score. Tone is part of this — does it sound like the examples in the Tone & Style Guide, or like a generic LLM customer service voice?

Reasoning quality — when the agent's reasoning shows it weighed the right rules, that matters more than the answer. When it's clearly bluffing past ambiguity, that's a problem.

What to put in your write-up

- Headline strict accuracy and per-field breakdown.
- At least 3 cases where you disagreed with the benchmark, with your reasoning. (See 'ambiguous_cases' in the benchmark JSON for hints — but don't limit yourself to those.)
- At least 2 cases where the agent failed and you can articulate why (prompt issue? missing context? model limitation?).
- A note on tone — was the agent's draft voice consistent with the tone guide, or did it drift?

How we'll evaluate your evaluation

We are explicitly **not** just looking at the accuracy number. A candidate scoring 75% with sharp self-critique and clear reasoning is more interesting to us than one scoring 95% with no reflection. We've seeded a few

benchmarks that are genuinely debatable — we want to see if you spot them and argue back. Building agents that work in production means knowing when the spec is wrong.

Things we'll notice (positively)

- A clear separation between rule-following (SOP) and judgement (escalation, tone).
- An agent design that handles all 20 messages with the same prompts — no per-message hacks.
- A grounded answer to 'what would you build next' — not just 'I'd add evals' or 'I'd add a vector DB'.
- Spotting at least one place where the SOP, the catalogue, or the tone guide contradict each other.

Things we'll notice (negatively)

- Hand-tuning prompts to pass specific benchmark cases.
- Draft replies that are obviously generic LLM voice — apologising five times, never referencing the actual problem.
- Quoting prices the catalogue marks as 'from'.
- Architectural complexity for its own sake — five agents and a router for a problem that fits in one prompt.